

# **Data Collection & Preprocessing Report**

(Submitted by: Ahsan Zafar H. Syed, NUID: 002801441)

## **1. Dataset Sources and Total Size**

For this project, three publicly available text datasets from the HuggingFace Hub were used. To meet the requirement of preparing a dataset of at least 1GB while preserving storage efficiency, all large datasets were loaded in streaming mode, preventing full downloads (which would otherwise exceed 40GB+).

The final dataset contains approximately 1GB of cleaned raw text sampled from:

### **Datasets Used**

1. Wikimedia/Wikipedia (English Dump)  
Provides factual, encyclopedic, well-structured text covering diverse topics.
2. HuggingFaceFW/FineWeb-Edu (sample-10BT subset)  
A curated subset of large-scale web crawls with general web-style content, including articles, blog posts, and discussions.
3. SetFit/AG News  
A smaller news dataset used to introduce journalistic writing styles for additional domain variety.

### **Why These Datasets?**

- Combines encyclopedic, web, and news domains
- Provides stylistic and linguistic diversity
- Supported by streaming, avoiding multi-GB downloads
- Stable, widely used datasets suitable for language model pretraining
- Sampling continued until the global target (~1.5GB raw cleaned text) was reached.

## **2. Cleaning Strategies and Reasoning**

High-quality preprocessing is essential for stable and efficient language model pretraining. The following pipeline was applied.

### **Cleaning Steps**

1. Lowercasing  
Ensures vocabulary consistency and prevents BPE token explosions caused by inconsistent casing.
2. HTML/Markup Removal  
FineWeb data often contains HTML fragments. A lightweight regex was used to remove <tags> and inline markup.

### 3. Whitespace Normalization

All repeated whitespace was collapsed into a single space using \s+ regex.

### 4. Length Filtering ( $\geq 50$ words)

Removes extremely short or low-information content (titles, ads, broken lines), keeping only meaningful text.

### 5. Deduplication (MD5 Hashing)

To ensure uniqueness, MD5 hashes were computed from the first 2000 characters of each cleaned document. Documents with identical hashes were discarded.

## Benefits of Cleaning

- Reduces noise and junk text
- Prevents the model from memorizing duplicates
- Improves the representativeness of token distributions
- Standardizes input format across heterogeneous sources

## 3. Tokenization Choices

Tokenization prepares raw text into tokenIDs for model consumption. The following choices were made:

### Tokenizer

- GPT-2 Byte Pair Encoding (BPE)  
Loaded using AutoTokenizer from Hugging Face.

### Vocabulary

- ~50,000 subword units  
(GPT-2 standard vocabulary size)

### Special Token Decisions

- GPT-2 has no built-in PAD token
- EOS token was reused as the padding token to allow batch padding

### Chunking Strategy (Block Size)

- Block size: 512 tokens
- All tokenized sequences were split into fixed-size blocks of max 512 tokens.

### Why Block Size = 512?

- Memory-efficient during training
- Commonly used for mid-size transformer training
- Allows good context range while keeping batch sizes manageable

## 4. DataLoader & Implementation Details

A fully custom PyTorch dataloader was implemented to handle batching, padding, and attention mask generation.

### Components

#### TokenBlockDataset

- Stores all token blocks as Python lists of token IDs.
- Implements standard `__len__` and `__getitem__`.

#### Custom collate\_fn

Handles:

- Dynamic padding to match the longest sequence in the batch
- Creation of the `attention_mask`
  1. 1 for real tokens
  2. 0 for padding tokens

#### PyTorch DataLoader

- `batch_size = 8`
- `shuffle = True`
- `collate_fn` injected via lambda
- Efficient batch assembly for training

### Saved Output

A sample batch containing:

- `input_ids`
- `attention_mask`

was saved to `sample_dataset.pt` as required.

## 5. Challenges Encountered

### 1. Deprecated Dataset Scripts

The older "wikipedia" dataset is no longer supported in Hugging Face.

Solution: use `wikimedia/wikipedia` (Parquet-based, streaming-compatible).

### 2. Memory and Disk Constraints

Loading Wikipedia and FineWeb normally would require tens of gigabytes.

Solution: streaming mode so only small portions are kept in memory.

### 3. Deduplication Efficiency

Deduping millions of documents requires storing hashes.

Solution: hash only first 2000 characters → drastically reduces memory footprint.

#### **4. Inconsistent Column Names**

Different datasets use different field names:

Wikipedia: "text"

#### **5. FineWeb: "text"**

AG News: "text" or "description"

Solution: dynamic field detection.

#### **6. Variable Sequence Length**

Batches contained variable-length sequences.

Solution: custom collate function to handle padding and attention mask generation.

### **6. Reflections on Impact to Pretraining Quality**

Preprocessing quality has an outsized impact on the performance, stability, and sample-efficiency of language model pretraining.

#### **Key Reflections**

##### **1. Cleaner input = more stable training**

Noise, HTML, duplicates, and broken text harm convergence. Cleaning improves token distribution and reduces gradient noise.

##### **2. Deduplication prevents memorization**

Language models easily memorize repeated documents. Deduplication ensures more efficient learning and better generalization.

##### **3. Domain-balanced sampling improves robustness**

Mixing encyclopedic, web, and news text exposes the model to diverse syntax and styles.

##### **4. Chunking at 512 tokens balances memory and context**

Larger blocks improve context understanding but reduce batch size; 512 is a good compromise.

##### **5. Tokenizer compatibility is essential**

Choosing GPT-2's tokenizer ensures a stable, well-tested subword vocabulary.

### **Overall Conclusion**

High-quality preprocessing is arguably as important as model architecture.

Clean, deduplicated, diverse data directly improves:

- Convergence speed
- Representation quality
- Downstream fine-tuning performance
- Generalization & robustness